

In this problem set you'll practice fitting and evaluating different predictive models to see if you can detect overfitting. You can load in the data via a link on Ed. For the following questions, fill in the space below with the R code you used.

1. Visualize the relationship between x and y with a `ggplot` and comment on what you see, describing the strength, direction, and shape/form of the association.
2. Split your data into training and testing sets; seventy percent of the data should be allocated to the training set.
3. Fill out the below table. You will fit a model with a polynomial having the degree specified, and report the testing and training RMSE in each case.

Degree	Training RMSE	Testing RMSE
1		
2		
3		
4		
5		
10		
20		
25		

To help you, here is some code that will calculate training RMSE for you, provided you have fit a linear model called `m1` and make a training set called `train`:

```
train |>
  mutate(yhat = predict(object = m1, newdata = _____),
         resid = _____) |>
  summarise(MSE = mean(resid^2))
```

You will need to modify this code slightly to help you find the testing RMSE. Write the modified code in the space below given a linear model `m1`.

4. Describe the pattern in the results you see and explain it.